


1. The concurrent execution of these two processes can result in one or both being blocked forever due to a deadlock situation. A deadlock occurs when each process holds a semaphore that the other process needs and vice versa. In this case, a deadlock can happen if the following sequence of events occurs:
  - a. foo acquires semaphore S.
  - b. bar acquires semaphore R.
  - c. foo tries to acquire semaphore R but foo is blocked because bar holds R.
  - d. bar tries to acquire semaphore S but bar is blocked because foo holds S.
 As a result, both processes are waiting for the semaphore held by the other which results in a deadlock, and neither can proceed.


2.

3.


a. **FCFS:**



**Non-Preemptive Priority (Smaller Priority Number Implies Higher Priority):**



**Round Robin with Quantum 30 ms:**



b. **FCFS:**

Avg. Waiting Time =  $(0 + 30 + 30 + 110) / 4 = 42.5$  ms

**Non-Preemptive Priority (Smaller Priority Number Implies Higher Priority):**

Avg. Waiting Time =  $(0 + 30 + 70 + 10) / 4 = 27.5$  ms

**Round Robin with Quantum 30 ms:**

Avg. Waiting Time =  $(20 + 10 + 70 + 70) / 4 = 42.5$  ms

c. **FCFS:**

Avg. Turnaround Time =  $(50 + 50 + 130 + 150) / 4 = 95$  ms

**Non-Preemptive Priority (Smaller Priority Number Implies Higher Priority):**

Avg. Waiting Time =  $(50 + 50 + 170 + 50) / 4 = 80$  ms

**Round Robin with Quantum 30 ms:**

Avg. Waiting Time =  $(70 + 30 + 170 + 110) / 4 = 95$  ms

4.

- a. The process can be scheduled and executed twice in a single round of the scheduler. This essentially allows the process to get a longer time slice or more CPU time in a single quantum compared to other processes.
- b. The major advantage of this scheme is that it allows a specific process to have more consecutive time slices without being preempted. This can be beneficial in scenarios where a process requires a larger continuous chunk of CPU time, and frequent preemption might result in less efficient use of the CPU.
- c. We can allocate longer quantum time for the high-priority processes

5.

- a. Available = Total - Sum of Allocation = 12 12 8 10 - 9 9 6 9 = 3 3 2 1
- b. CurrentNeed = Max - Allocation

Process	Allocation	Max	CurrentNeed
	A B C D	A B C D	A B C D
<b>P0</b>	2 0 0 1	4 2 1 2	2 2 1 1
<b>P1</b>	3 1 2 1	5 2 5 2	2 1 3 1
<b>P2</b>	2 1 0 3	2 3 1 6	0 2 1 3
<b>P3</b>	1 3 1 2	1 4 2 4	0 1 1 2
<b>P4</b>	1 4 3 2	3 6 6 5	2 2 3 3

c.

Process	Allocation	Available	CurrentNeed
	A B C D	A B C D	A B C D
P0	2 0 0 1	3 3 2 1	2 2 1 1
P4	3 1 2 1	5 3 2 2	2 1 3 1
P2	2 1 0 3	5 3 2 2	0 2 1 3
P3	1 3 1 2	5 3 2 2	0 1 1 2
P4	1 4 3 2	6 6 3 4	2 2 3 3
P1	3 1 2 1	7 10 6 6	2 1 3 1
P2	2 1 0 3	10 11 8 7	0 2 1 3

**Safe State:** P0, P3, P4, P1, P2

d.

Request (1, 1, 0, 0) <= Available (3 3 2 1) **TRUE**

Request (1, 1, 0, 0) <= Need (2, 1, 3, 1) **TRUE**

Need(i) = Need(i) - Request(i) = 2 1 3 1 - 1 1 0 0 = 1 0 3 1

Available = Available - Request(i) = 3 3 2 1 - 1 1 0 0 = 2 2 2 1

Allocation(i) = Allocation(i) + Request(i) = 3 1 2 1 + 1 1 0 0 = 4 2 2 1

Process	Allocation	Available	CurrentNeed
	A B C D	A B C D	A B C D
P0	2 0 0 1	2 2 2 1	2 2 1 1
P4	4 2 2 1	4 2 2 2	1 0 3 1
P2	2 1 0 3	4 2 2 2	0 2 1 3
P3	1 3 1 2	4 2 2 2	0 1 1 2
P4	1 4 3 2	5 5 3 4	2 2 3 3
P1	4 2 2 1	6 9 6 6	1 0 3 1
P2	2 1 0 3	10 11 8 7	0 2 1 3

**Safe State:** P0, P3, P4, P1, P2

e. Request (0, 0, 2, 0) <= Available (3 3 2 1) **TRUE**

Request (0, 0, 2, 0) <= Need (2, 1, 3, 1) **TRUE**

Need(i) = Need(i) - Request(i) = 2 2 3 3 - 0 0 2 0 = 2 2 1 3

Available = Available - Request(i) = 3 3 2 1 - 0 0 2 0 = 3 3 0 1

Allocation(i) = Allocation(i) + Request(i) = 1 4 3 2 + 0 0 2 0 = 1 4 5 2

Process	Allocation	Available	CurrentNeed
	A B C D	A B C D	A B C D
<b>P0</b>	<del>2 0 0 1</del>	<del>3 3 0 1</del>	<del>2 2 1 1</del>
<b>P1</b>	<del>4 2 2 1</del>	<del>3 3 0 1</del>	<del>1 0 3 1</del>
<b>P2</b>	<del>2 1 0 3</del>	<del>3 3 0 1</del>	<del>0 2 1 3</del>
<b>P3</b>	<del>1 3 1 2</del>	<del>3 3 0 1</del>	<del>0 1 1 2</del>
<b>P4</b>	<del>1 4 5 2</del>	<del>3 3 0 1</del>	<del>2 2 1 3</del>

**Not Safe**